

KEDAR LIMBALKAR

kedarlimbalkar@gmail.com | +91-9284374129 | Bengaluru, Karnataka | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Machine Learning Engineer and Data Scientist with hands-on expertise in deep learning, NLP, LLMs, computer vision, and end-to-end ML pipelines. Proficient in RAG workflows, model fine-tuning, prompt engineering, MLOps, and Docker deployment. Proven track record of building and deploying production-grade AI systems with measurable accuracy and latency improvements.

TECHNICAL SKILLS

Languages & Libraries: Python (Pandas, NumPy, Scikit-learn, XGBoost, Keras, PySpark, Matplotlib, Seaborn), SQL

ML / Deep Learning: TensorFlow, CNN (Convolutional Neural Networks), Computer Vision, Feature Engineering, Statistical Analysis.

Generative AI & LLMs: LangChain, LangGraph, FAISS, Llama 3.3, Hugging Face, Prompt Engineering, Model Fine-Tuning, RAG (Retrieval-Augmented Generation), Vector Embeddings, Transformer Models, Semantic Search, Multi-Agent Systems

NLP: Natural Language Processing, Text Classification, Sentiment Analysis, Named Entity Recognition, Tokenization, Text Preprocessing, Word Embeddings

MLOps & DevOps: Docker, Git, GitHub, CI/CD, MLOps Deployment, Streamlit, OpenCV, REST API, Apache Airflow

Data Engineering: ETL Pipelines, Data Modeling, Data Wrangling, Data Visualization, EDA (Exploratory Data Analysis)

Databases & Cloud: MySQL, BigQuery, Snowflake

WORK EXPERIENCE

Intucate | Business Analyst Intern

January 2026 – April 2026

- Analyzed content asset performance metrics using data analysis and EDA techniques, identifying trends that improved educational engagement outcomes by 15% for target demographics.
- Collaborated with product and engineering teams to translate business requirements into structured workflows, aligning roadmap priorities across 3+ departments.

PROJECTS

Conversational AI for Retail | [GitHub](#) | [Drive](#)

January 2026 – February 2026

Technologies: Python, LangGraph, LangChain, Llama 3.3, FAISS, RAG, Vector Embeddings, NLP, Prompt Engineering, Multi-Agent Systems

- Preprocessed 44,884 retail records using a 3,000/500 token chunking strategy to build a FAISS vector knowledge base, achieving 100% semantic integrity for rapid retrieval.
- Designed a multi-agent LangGraph RAG (Retrieval-Augmented Generation) pipeline achieving 92% domain accuracy and sub-1.2s query latency, reducing manual customer service workload by 40%.
- Applied prompt engineering and model fine-tuning with Llama 3.3, reducing hallucination rate by 18% and improving response precision across 10,000+ user interactions.

Emotion Recognition System | [GitHub](#) | [Drive](#)

February 2026 – March 2026

Technologies: Python, Keras, TensorFlow, CNN, OpenCV, Streamlit, Computer Vision, Deep Learning, Transfer Learning, Batch Normalization

- Built a CNN (Convolutional Neural Network) in Keras to classify 7 emotional states from 35,000+ facial images, achieving 88% test accuracy using Dropout and Batch Normalization layers.
- Deployed a real-time computer vision inference pipeline with OpenCV and Streamlit, delivering 30+ FPS; applied transfer learning and hyperparameter tuning to reduce overfitting by 20%.
- Evaluated model performance using cross-validation and confusion matrix analysis, improving precision by 12% over the baseline CNN architecture across all 7 emotion classes.

EDUCATION

MS in Computer Science: Machine Learning & AI Engineering — Woolf University

2025 – 2027

Relevant Coursework: Machine Learning, Artificial Intelligence, Deep Learning, NLP, Data Science, Statistical Modeling

B.Tech in Computer Engineering — Government College of Engineering, Jalgaon

2021 – 2025

Relevant Coursework: Data Structures, Algorithms, Statistics, Databases, Computer Vision, Software Engineering

CERTIFICATIONS & ACHIEVEMENTS

Full Stack Data Science and AI | AlmaBetter | [View Certificate](#)

- Completed 6 production-ready ML/DL projects across 8+ datasets covering machine learning, deep learning, NLP, statistical modeling, data analysis, and Docker-based MLOps deployment.
- District Winner — National Level Computing Intelligence Olympiad (CIO): ranked 1st at district level in logical reasoning and data interpretation among 500+ participants.
- Core Member, Technical Council 2025 — coordinated 3 national-level technical events, managing logistics and cross-team communication for 200+ attendees.